



The Full-Stack GTM Roadmap

A disproportionate amount of value in Go-to-Market is concentrating among professionals who evolve with AI — expanding across the full funnel, enabled by agents, agentic systems, and expert command of modern tool architectures. This is the playbook for becoming one of them.

V1.1 — CLAY GTM ENGINEERING BLOG INTEGRATED

ALEX · 2026

How to Read This Roadmap

This document is V1.1 — meant to be **argued with, not accepted**. Two parts that stand alone and serve different purposes: Part 1 defines the skill domains required to become the Full-Stack GTM Pro archetype, and Part 2 runs the job hunt itself as an AI-native full-funnel GTM motion.

Part 1 — Coursework

Nine capability domains organized by depth target, with core concepts, canonical reading, and hands-on benchmarks. Each domain maps to one or more rungs of the Clay maturity model:

Data Foundation → Data Modeling → Data Activation.

Part 2 — Job-Hunt Funnel

The search run as an AI-native GTM motion — from market mapping through offer close and ramp. Each stage has objective, actions, tools, artifacts, metrics, and a "what the top 1% does" differentiator.

What Changed in V1.1

- Consultative-seller thesis added as positioning bullseye
- Three-rung maturity model as operating sequence
- GTM alpha concept + org-model taxonomy with named companies
- Six portfolio play patterns from real Clay builds
- Tier-1 target list: Intercom, Canva, Notion, Anthropic, Ramp, Verkada, Rippling
- Clay's three-part assessment rubric integrated into Stage 6

① Source material: Five Clay JDs (Forward Deployed GTM Engineer, Head of GTM Architecture, Product Marketing, Analytics Engineer, Enterprise Growth Strategist-CS), the Clay GTM Engineering blog, the gtm-agents marketplace (243 skills across 67 plugins), and the Empire State Events Pipeline architecture.

Part 0 — The Archetype You're Optimizing For

Before the coursework, a grounding on what the five JDs and the Clay blog collapse into. This is the target — the role archetype you are training toward.

The Thesis, Stated Plainly

Clay's GTM engineering post: *"Just as Harvey uses lawyers to sell legal software, Clay uses GTMEs to sell Clay."* You are not an AE learning to code — you are the consultative seller who's **lived the workflow AND can now build against it**. That is Clay's explicit hiring mental model.

The Operating Sequence

Rung 1 — Data Foundation: Clean, deduped CRM + warehouse records. **Rung 2 — Data Modeling:** Propensity scores, ICP attributes, AI-data points that predict purchase. **Rung 3 — Data Activation:** Revenue-generating workflows — automated outreach, churn-risk nudges, sales enablement. "Most companies stumble on the first rung."

GTM Alpha — The Concept You're Selling Yourself As

GTM tactics have commoditized. The edge lives in **unique data + differentiated plays** executed in hours, not quarters. Your narrative — interviews, LinkedIn, portfolio — should demonstrate signal density, narrow segmentation, and execution speed. The opposite is "hire 10 more reps for every 10% pipeline lift."

Where the Function Lives: Two Org Models

Model 1 — GTME Inside RevOps

Intercom — Alexander DeMoulin's GTM Ops pilots plays; GTM Systems scales to production. **Canva** — Robert Jones's GTM AI team automates workflows. **Notion** — Three teams: Evan Peters (RevOps + Strategy → CRO), Jalal Iftikhar + Shelley Lee (BizTech → CFO), Theo Bleier (AI Engineer → CTO). **Anthropic** — GTM Infrastructure under Adam Wall; "Productivity Engineering" squad wields Claude + Clay.

Model 2 — GTME Inside Growth

Ramp — Growth Platform squad under Keyan Sarrafzadeh; two-week sprints. **Verkada** — GTMEs under Davide Grieco in Growth; 80-100 meetings/rep/month via SDR automation. **Rippling** — GTM engineering inside Noah Adelstein's international + US growth teams; automated outbound + direct mail.

📌 Cross-model lessons: (1) Start in ops, federate later. (2) Balance rapid prototyping with stable implementation. (3) Data hygiene is non-negotiable — every win traces back to clean, enriched, well-modeled data.

Part 1 — Coursework: Structure & Depth Targets

Nine capability domains, each with a "why it matters" tied to the JDs, core concepts, canonical reading, and hands-on benchmarks. Depth targets define where you must land — the realistic goal for an AI-native GTM senior hire is **Own in 3–4 domains, Do in the rest, Recognize nothing**.

 **Recognize**
Minimum bar. You can read about it, follow a conversation, and know when to ask for help. Not acceptable as a terminal state for a senior hire.

 **Do**
Working bar. You can execute competently with reference material, produce work product, and talk through trade-offs. This is the floor everywhere.

 **Own**
Senior bar. You can design systems, teach others, and be the senior voice in the room. Target for 3–4 domains — your commercial, GTM systems, GTM engineering, and AI/agent domains.

The Nine Domains at a Glance

#	Domain	Depth Target	Rung Mapping	Why It Matters
1	Commercial & Enterprise Sales	Own	All three	Foundation the other eight domains sit on
2	Full-Funnel GTM Systems	Own	All three	Head of GTM Architecture axis; highest-leverage differentiator
3	GTM Engineering Craft	Own	All three	Forward Deployed GTME JD; fastest proof surface
4	Data & Analytics Fluency	Do	Rungs 1 + 2	Earn the right to be on rungs 2 and 3
5	AI / Agent Engineering	Own	Rungs 2 + 3	AI multiplier proof; Empire State pipeline is the artifact
6	Product Marketing & Narrative	Do → Own	Rungs 2 + 3	Documentarian angle is native PMM; gap is discipline
7	Customer Success & Expansion	Do	Rungs 2 + 3	Enterprise Growth Strategist axis; land-and-expand motion
8	Leadership & Org Design	Do	All three	Head of GTM Architecture JD; 0→1 builder DNA
9	Written Communication	Own	All three	Weapon, not afterthought; ships documents not just decks

Domain 1 — Commercial & Enterprise Sales Mastery

DEPTH TARGET: OWN

YOUR MOAT — SHARPEN, DON'T JUST MAINTAIN

Every JD assumes commercial seriousness. Your selling DNA is the foundation the other eight domains sit on — but in an AI-native context, the bar has moved from "quota-carrying" to "**quota-carrying + operates like a revenue architect.**"

Core Concepts

- Modern enterprise deal shape: committee-based buying, procurement + security gates, multi-threading, economic vs. technical vs. champion personas
- Value-based selling: quantified business outcomes vs. feature parity
- Qualification frameworks as **inspection tools**, not scripts

Frameworks & Canonical Sources

- **MEDDPICC** — lingua franca of enterprise pursuits; already in your gtm-agents repo
- **Command of the Message / Command of the Sale** (Force Management) — value-first messaging at the deal level
- **Winning by Design** — bowtie model, SPICED discovery
- **The Challenger Sale** — teach, tailor, take control
- **Never Split the Difference** (Voss) — late-stage and offer-stage leverage

Hands-On Benchmarks

- Produce a written **deal strategy memo** (3–5 pages) on a real prospect using MEDDPICC + SPICED, including a multi-thread org map and risks-and-mitigations register
- Record yourself running a mock discovery call with commercial-insight framing; self-review against SPICED and MEDDPICC
- Draft an **exec-sponsor email** that earns a 30-min meeting with a CRO-level buyer using only public signals

Concept: Value-Based Selling

It is not "talk about benefits instead of features." It is arithmetic. Build a model with the prospect: (a) current-state cost of the problem, (b) probable-state cost with your solution, (c) the delta, and (d) the payback period. If you cannot quantify the delta, you cannot defend the price at procurement.

Domain 2 — Full-Funnel GTM Systems Thinking

DEPTH TARGET: OWN	HIGHEST-LEVERAGE DIFFERENTIATOR FOR A SENIOR IC
-------------------	---

Clay's Head of GTM Architecture JD: *"Architect the GTM operating rhythm end-to-end. Design inspection and accountability cadences. Operationalize comp rollout."* This is the move from **"I hit quota"** to **"I understand why the quota is what it is, and I can redesign the engine that produces it."**



The Bowtie Funnel

Acquisition → retention → expansion. SaaS maturity makes the **right side of the bowtie** the value center. Net retention and expansion revenue are the metrics that matter most to a CFO.



Capacity Modeling

Reps × productivity × ramp × territory coverage → quota → capacity gap. Build a **20-AE team model** hitting \$40M new ARR with ramp curves, productivity assumptions, attrition, and a hiring plan. Stress-test it.

Canonical Sources

- **The Revenue Architecture** — Winning by Design (installed as skill)
- **From Impossible to Inevitable** — Ross & Cannon
- **The SaaS CFO blog** — Ben Murray's operator-language metrics
- **Tomasz Tunguz, Jason Lemkin, SaaStr** — public benchmarks for every funnel metric
- **gtm-agents plugins** — revenue-analytics, revenue-forecasting-pipeline, enterprise-sales, sales-pipeline, sales-operations



Operating Rhythm

Forecast calls, pipeline councils, QBRs, forecast-to-actual variance reviews, rules of engagement. A junior operator asks "are we hitting plan?" A senior operator asks *"do we have the meetings, artifacts, and inspection loops that would let us know if we weren't?"*



Territory & Comp Design

Accelerators, SPIFFs, multipliers, clawbacks, the pathology of "paying for activity." Carve-outs, account assignment, deal-split rules, lead-to-opp routing. Draft a **rules-of-engagement document** for a 2-segment sales org.

Hands-On Benchmarks

- Build a **capacity model in Excel/Sheets** for a 20-AE team hitting \$40M new ARR; stress-test with attrition and ramp scenarios
- Draft a **"rules of engagement" document** for Commercial (≤500 employees) vs. Enterprise (>500) segments
- Write a **one-page operating rhythm proposal** for a Series B AI-native company: which meetings exist, what each decides, who owns each, what artifacts come out
- Run a **funnel decomposition retrospective** on one of your own closed deals: which stage drove velocity, which created slippage

Domain 3 — GTM Engineering Craft

DEPTH TARGET: OWN

RUNGS 1 + 2 + 3 — WHERE THE MATURITY MODEL BECOMES EXECUTABLE

This is the Forward Deployed GTM Engineer JD near-verbatim and the backbone of Clay's entire blog thesis. GTMEs are *"an internal product team that serves the GTM organization — they identify problems, write specs, ship prototypes, and scale what works."* Every build should lead with two numbers: **meetings booked** and **hours saved**.

Core Concepts

- **Workflow orchestration:** event-driven automation, webhooks, triggers, conditional branches, retries, idempotency
- **Enrichment waterfalls:** cost-aware provider sequencing, credit budgeting, quality scoring, suppression
- **Identity resolution:** deduping across sources, match keys, source-of-truth watermarking
- **Signal architecture:** intent, product usage, web, job-postings, news, funding, technographic — classified, weighted, decayed, routed
- **API literacy:** REST, OAuth, rate limits, pagination, idempotency keys, retries with backoff

Tools to Actually Learn

- **Clay** — build 5 workflows end-to-end: inbound enrichment, outbound list-build, signal-based trigger, CRM-write with dedup, Slack alert
- **HubSpot** — workflows, custom properties, associations, custom objects, reports
- **Salesforce Admin fundamentals** — Trailhead Admin Beginner (~10 hours)
- **n8n, Zapier, Make** — you have n8n chops; lean in
- **Postman / curl** — API sandboxing; every GTM engineer lives here

Portfolio Play Patterns (from Real Clay Builds)

- **Inbound scoring + auto-routing** — scores signups, auto-assigns to best-fit rep (Manny & Osman: \$25k+ potential)
- **Transcript → CRM backfill** — LLM extracts firmographics from call recordings, writes to Salesforce (Blake)
- **Signal digest in Slack** — account-awareness signal piped to sellers for call prep (Osman)
- **VIP event-invite cross-reference** — matches attendee list to CRM + geolocation (Spencer)
- **Programmatic partner sourcing** — scrapes social + website profiles, scores fit (Tommy)
- **Reverse-IP dynamic outreach** — triggers personalized email while interest is hot (Pranav)
- **Ticket-spike → docs drafting** — flags support surges, drafts help-center articles (Bruno)
- **Expansion radar** — scans support threads for enterprise-feature requests before renewal (Osman)

Three Capstone Builds

1	2	3
Capstone 1 — Clay-Equivalent in Your Own Stack Build in n8n + HubSpot + Apollo/Clay: ingest contact → enrich with 2-provider waterfall → classify fit via LLM → route to sequence → log to Notion audit table. Report: hours saved vs. manual baseline.	Capstone 2 — Event-Driven Outbound Engine Ingest a signal source (job-postings, BuiltWith, funding), filter by ICP rules, generate personalized first-touch via LLM, write to CRM, trigger sequence. End-to-end in ≤48 hours. Report: meetings booked per 100 contacts.	Capstone 3 — Forward-Deployed Consulting Simulation Pick a Tier 1 target company. Write a 5-page implementation proposal: current-state diagnosis, architecture diagram, 6-week build plan, success metrics. Share on LinkedIn — portfolio artifact and job-search asset.

Domain 4 — Data & Analytics Fluency

DEPTH TARGET: DO

RUNGS 1 + 2 — EARN THE RIGHT TO BE ON RUNGS 2 AND 3

Clay's blog is explicit: *"Data hygiene is non-negotiable, and every win traces back to clean, enriched, well-modeled data."* You will not master the full modern data stack — but you need to understand what each layer does, speak SQL to a working level, and be able to look at a dbt model and understand it. **"Most companies stumble on the first rung."** A fluent GTM engineer recognizes that and fixes it before attempting fancy activation.

Core Concepts

- **SQL at working depth:** joins, window functions, CTEs, subqueries, aggregations, date math, NULL handling
- **Modern data stack layers:** sources → ELT (Fivetran, Segment) → warehouse (Snowflake, BigQuery) → transformation (dbt) → orchestration (Dagster) → BI (Hex, Sigma) → reverse ETL (Census)
- **Dimensional modeling:** facts vs. dimensions, star schemas, slowly-changing dimensions, grain
- **Semantic layer / metrics layer:** why having "ARR" defined in one place matters
- **Experiment design:** hypothesis, control/treatment, sample sizing, guardrails, p-values
- **Funnel math & cohort analysis:** stage-to-stage conversion, velocity, retention curves, cohort decay

Canonical Sources

- **Mode's SQL tutorial** (free) — gold standard for working-SQL fluency
- **dbt Fundamentals** (free, ~4 hours) — get to Recognize; build one project to reach Do
- **Designing Data-Intensive Applications** — Kleppmann; skim data modeling chapters
- **gtm-agents plugins:** analytics-pipeline-orchestration, business-intelligence, revenue-analytics, marketing-analytics

Hands-On Benchmarks

- **Cohort retention analysis from a CSV** — given signup dates and monthly usage, produce a cohort retention heatmap with a 1-paragraph narrative
- **Build one dbt model** — a stg_ model and a fct_ model transforming a raw CRM export into an "opportunities" fact table
- **Metrics dictionary** — one-page definitions of ARR, NRR, GRR, CAC, Payback, Magic Number, Rule of 40 with testable formulas
- **Experiment post-mortem** — pick a published A/B test; re-derive whether the result was significant given the sample size

Concept: The Semantic Layer

In a pre-semantic-layer world, "ARR" is defined differently in 11 dashboards and 4 executive reports — the Head of Sales and the CFO report different numbers, and neither can be reconciled in under 3 hours. In a post-semantic-layer world (dbt metrics, Cube, LookML), ARR is defined *once*, in code, with tests and version history. This is why AI-native companies care about dbt: **agent-driven analytics is only safe if the metrics are bulletproof.**

Domain 5 — AI / Agent Engineering

DEPTH TARGET: OWN

YOUR "AI MULTIPLIER" PROOF — YOU'RE ALREADY 60-70% THERE

You are being hired *because of* your AI fluency, not despite it. The Empire State Events pipeline is the proof artifact for this domain. The gap to close: from **"builds useful AI workflows"** to **"speaks the engineering vocabulary of LLMs, agents, and evals fluently enough to sit at a product-engineering table."**

Core Concepts

- **LLM fundamentals:** tokens, context windows, temperature, top-p, structured outputs, function/tool calling, streaming
- **Prompt engineering craft:** system prompts, few-shot vs. zero-shot, chain-of-thought, prompt chaining, XML tags as structure
- **RAG:** embeddings, vector stores, chunking strategies, hybrid retrieval, reranking, grounding
- **Agents:** LLM as planner-executor with tools; loop shapes, planning vs. reactive agents, multi-agent systems
- **Context engineering:** the discipline of what goes into the context window, when, and why — the successor to "prompt engineering"
- **Evals:** offline evals, golden datasets, LLM-as-judge, regression tests for prompts — without evals you are building on sand
- **Cost + latency engineering:** model selection (Haiku/Sonnet/Opus), caching, parallelism, batching

Seminal Sources (Pick ≥6)

- **"Attention Is All You Need"** — Vaswani et al., 2017; read the abstract, skim the rest
- **"Language Models are Few-Shot Learners"** — GPT-3 paper, 2020
- **"ReAct: Synergizing Reasoning and Acting"** — Yao et al., 2022 (the agent-loop paper)
- **Anthropic's "Building Effective Agents"** (2024) — cleanest taxonomy of agent patterns in public writing
- **Anthropic's MCP documentation** — read the spec once to level up from user to fluent user
- **Simon Willison's blog** — ongoing practitioner commentary, high signal
- **Lilian Weng's blog** (OpenAI) — synthesis posts on agents, hallucination, evals

Hands-On Benchmarks

- **Build one non-Empire-State agentic system end-to-end** — 5-10 hour build, single LLM + 3-5 tools + 1 eval suite (e.g., inbound-lead-qualifier agent, competitive-brief-agent)
- **Write a prompt eval harness** for the Empire State event-research skill — 10 golden examples, LLM-as-judge rubric, score + regression log
- **Publish one technical write-up** — blog, LinkedIn long-form, or GitHub README detailed enough that another operator could rebuild it
- **Ship one MCP server** — use the mcp-builder skill; even a trivial one (e.g., reads your Notion Events DB) is a resume credential

Concept: Context Engineering

Two years ago, the unit of work was "write a good prompt." Today, it is *"design what goes into the context window, when, and why"* — because modern agents chain many LLM calls, pull from many sources, and compose many sub-prompts. Do you retrieve and inject, or tool-call and let the model pull? Do you summarize prior turns or keep them raw? Do you use an orchestrator model (Opus) to plan and worker models (Haiku) to execute? **This is engineering, not copywriting.** If you have an opinion on this shift and can defend it, you're a senior.

Domain 6 — Product Marketing & Narrative

DEPTH TARGET: DO → OWN

HEADED TOWARD OWN AS YOU BUILD A PUBLIC BODY OF WORK

The modern AI-native PMM: *"Own Clay's core narrative and messaging frameworks. Develop segment and persona-specific positioning. Build and maintain competitive intelligence."* Your documentarian angle from Empire State Events is a native PMM skill in disguise — you're already doing thought-leadership content. The gap is turning instinct into **discipline**.

Core Concepts

- **Positioning** (not messaging): who the product is for, against what alternatives, for what value. Positioning is strategy; messaging is expression.
- **Message architecture**: core promise, pillars, proof, audience-specific expressions
- **Jobs-to-be-Done (JTBD)**: what job is the customer hiring your product to do, in what circumstance
- **Competitive intelligence as a system**: battlecards, win/loss analysis, feature-parity vs. narrative-differentiation
- **Proof hierarchy**: logos, case studies, ROI models, analyst recognition, peer-review (G2, Gartner), benchmarks
- **Enablement**: turning positioning into sales reps' working vocabulary — where most PMM work fails
- **Pricing & packaging**: tiers, fences, value drivers, price-to-willingness-to-pay

Canonical Sources

- **April Dunford — "Obviously Awesome" + "Sales Pitch"** — the modern canon on positioning; read both
- **Andy Raskin's "Strategic Narrative" essays** — the 5-part deck structure every AI-native startup now uses
- **Geoffrey Moore — "Crossing the Chasm"** — still load-bearing for any segmented B2B motion
- **Emily Kramer (MKT1 newsletter)** — practitioner cadence on PMM + growth
- **Packy McCormick, Tomasz Tunguz, Stratechery** — narrative modeling at the company/category level

Hands-On Benchmarks

- **Write a positioning one-pager for yourself** using Dunford's template: competitive alternatives, unique attributes, value, best-fit customers, market category. Doubles as your job-hunt asset.
- **Build a competitive battlecard** for Clay vs. one alternative (Apollo, Outreach, HubSpot, Koala, Common Room). Share it publicly — PMM proof and a signal to Clay that you get them.
- **Write one "strategic narrative" piece** (Raskin-style 5-part) on the GTM-engineering category — ~1,500–2,000 words, pinned on LinkedIn.
- **Do a JTBD synthesis** of 3 people you've met at NYC AI events. What job were they hiring AI to do? Compare across them.

Concept: Command of the Message

Common confusion: sales methodology frameworks (Challenger, MEDDPICC, Command of the Message) are often thought of as PMM outputs. They aren't — they are *how the sales team uses* the positioning PMM produces. PMM owns the message house; sales enablement builds the call scripts; individual reps adapt it for prospects. If you can cross that handoff — write the positioning *and* operationalize it into the sales motion — you are unusually valuable, because **most PMMs can't do the second part and most AEs can't do the first**.

✔ **Progress Tracking:** Mark each benchmark as complete as you ship it. The three capstones in Domain 3, the eval harness in Domain 5, and the positioning one-pager in Domain 6 are your highest-signal portfolio artifacts — prioritize these first when sharing publicly.

Stage 6 — Evaluation: Teach to Clay's Rubric

Every interview at an AI-native GTM company is a live test. Clay's published assessment framework is the implicit rubric at Intercom, Notion, Ramp, Verkada, and Anthropic. Run every process to best-in-class — your reputation compounds regardless of outcome. The assessment has three parts:

1	2	3
Part A — Business Problem Investigation They hand you a fuzzy problem (e.g., "trial-to-paid won't budge"). Strong candidates ask the right first questions: Who is our ICP? What journey signals indicate readiness? What do winning customer conversations look like? Explore both market levers (intent data) and internal levers (SDR-to-AE routing). Specificity is the tell — not "I'd look at the funnel" but "I'd pull the last 30 trial cohorts, cut by ICP fit score, and compare 7/14/30-day activation against converted trials; my hypothesis is the gap is D3-D7."	Part B — Systems Sketch Turn the insight into a working flow. Pick unique data points to track (rung 2), keep data clean (rung 1 — don't skip this), plug into a Clay-equivalent workflow (rung 3), then test and iterate. The signal they want: <i>hypothesis</i> → <i>instrument</i> → <i>measure</i> → <i>iterate</i>, not <i>build</i> → <i>launch</i> → <i>hope</i>.	Part C — Mini Build Challenge Actual build, any tool. Design and validate data points that predict churn. What separates good from weak: explain assumptions and fallback logic, address suppression and multi-channel sequencing, state success criteria in advance, and ship a Loom walkthrough. The Loom is disproportionately effective — see Stage 5.

Clay's Published Red Flags — Memorize and Avoid

Tunnel Vision One solution, no alternatives considered.	Single-Channel Thinking Only email or only outbound — no multi-channel recognition.
Deliverability Blindness Outbound designed without domain warming, suppression, or fatigue management.	No Measurement Plan A build without stated success verification — self-disqualifying.

Stage 7 — Offer Close: Multi-Process Leverage + Negotiation

Never negotiate with one offer in hand. Time-box processes so offers land within 2 weeks of each other — this is the hardest operational part of job-hunting and the single most financially material move you'll make. Negotiate every component, not just base.

Compensation Components — Negotiate All

- Base salary
- Variable / OTE
- Sign-on bonus (offset equity cliff loss)
- Equity: amount, strike, refresh cadence, acceleration
- Accelerators and SPIFFs
- Promotion path + review cycle
- Quota structure (year 1 ramp, territory)
- Remote / hybrid terms
- Benefits: PTO, parental leave, equipment, learning budget

Negotiation Craft — Voss Playbook

- **Anchor high but defensibly** — ground in value math, not market data
- **Mirror + label emotions** to surface unstated objections
- **No-oriented questions:** "Would it be ridiculous to consider X?"
- **Let them solve your problem:** "How am I supposed to take a 20% base cut?"

Tools & Artifacts

- Comp-comparison model in Sheets — 2-hour build
- Levels.fyi + Pave.com for benchmarks
- Trusted senior mentor to pressure-test your ask

Metrics to Track

- Total comp delta vs. initial offer — target **10–25% uplift**
- Number of terms improved — not just base

📌 **What the top 1% does:** Brings a quantified, defensible case grounded in year-1 value delivery — not market comp data. "Here's the pipeline I'll build, here's the revenue I'll unlock, here's the comp that makes the math work." Market data is a backstop; value math is the ask. The acceptance letter is also your first day — send a "here's what I'll do in 60 days" note when you accept.

Stage 8 — Ramp: The First 90 Days as the Final Interview

By day 90, you are visibly the highest-performing hire the company has made this year. Treat the first 90 days as a public body of work and a team onboarding template at once.

1

Day 0–30: Observe + One Visible Win

Meet every stakeholder the role touches — AE team, ops, marketing, product, CS, finance. Ship one small, visible, well-executed thing in weeks 2–3: a workflow, a playbook, an analysis. Not the big strategic play — a credibility down-payment. Run your own 30-day diagnostic: where is the operating rhythm strong? Where weak? What would you change first?

3–5

Stakeholder meetings per week

Target in weeks 1–3

2

Day 31–60: Propose + Sequence

Present your diagnostic to your manager. Convert it into 2–3 medium-scoped initiatives for the next 60 days, prioritized. Start working cross-functionally — full-funnel fluency applies from day 1 of the job.

3

Shipped artifacts

By day 30, 60, and 90

3

Day 61–90: Ship the Strategic Play

Your proposed initiative ships in this window — visible, measurable, cross-functional. Request your first formal performance conversation. Set the year-1 narrative: what you'll own, what you won't, what you'll expand into. Document in public if culturally appropriate — your writeup becomes the next proof artifact.

90

Days to highest-performing hire

Manager already thinking about where you go
next

Iteration Passes: From V1.1 to Field-Ready

With the Clay blog integrated, four iteration passes remain. Each closes a specific gap between "structurally sound" and "field-tested."

1 Fidelity Pass on Part 1

Pressure-test domain priorities, depth targets, and the 12-week track design against your actual calendar and constraints. Drop or downshift what doesn't fit. Key open lever: weekly hour budget — see Confidence section.

3 Cross-Map Coursework onto Empire State Events Pipeline

Map each track's weekly work against what you're already building — event-research, content, pattern-synthesis, project-ideation. The Domain 3 play-pattern table is the connective tissue. Every capstone from Part 1 is built *on top of* Empire State, not next to it — this is where "real-world training program" earns its keep.

2 Fidelity Pass on Part 2


Red-team stage definitions, metrics, and differentiators. At least 2 stages have missing contingencies: Stage 5 outreach cadence and Stage 7 multi-process timing. Stress-test them.

4 Second-Order Build Targets from the Clay Blog

The blog names specific GTME team members at Clay — Manny, Osman, Spencer, Tommy, Blake, Bruno, Pranav. Their published plays are the highest-resolution public record of what "good" looks like. Pick 2–3, build your own version, publish with a walkthrough. Becomes both portfolio and a conversation-opener into Clay itself.

Confidence on V1.1: Where It's Solid and What's Still Open

Confidence scores reflect integration of Clay's published material. The single biggest open variable is your actual weekly hour budget.




The Critical Open Variable

The 12-week sprint assumes ~16–20 hours/week of directed learning + building while holding your current role. If that number is 8 instead of 18, the whole sequence compresses to 24 weeks and emphasis shifts toward Track A (Build) at the expense of Track B (Study). Lock this number before committing to the calendar.

What Moves Confidence Highest

- 2–3 practitioner conversations to field-test the role archetype
- Locking your weekly hour budget and stress-testing the 12-week calendar
- Red-teaming Stage 5 outreach cadence and Stage 7 multi-process timing
- Building 2–3 Clay team member play patterns into publishable artifacts



FULL-STACK GTM ENGINEER

DOMAINS 7-9

Full-Stack GTM Engineer: Part 2

This is the second half of a living progress-tracking presentation — a course load, project log, and public portfolio combined. The thesis: a disproportionate share of GTM value will concentrate in professionals who evolve with AI to work across the full funnel — enabled by agents, agentic systems, and expert use of modern tool architectures. This presentation tracks the domains, benchmarks, and job-hunt funnel to get there. Use it to track progress, share learnings, and signal your operating level to the right people.

Customer Strategy & Expansion

Enterprise Growth Strategist and Forward Deployed GTME roles resemble consulting more than traditional CSM — discovery, implementation, and value realization. Every enterprise deal now requires selling the expansion path at initial close. CS fluency raises your ceiling whether you target CS-adjacent or pure sales roles.

Core Concepts

- **Land-Adopt-Expand** — default model for modern enterprise SaaS; multi-year account strategy
- **Success planning** — mutual success plans, joint exec sponsors, milestones, KPIs, governance
- **QBR/EBR mechanics** — what the artifact says, how to run the meeting
- **Time-to-value** — the single most predictive metric for expansion
- **Customer health scoring** — quantitative (usage, engagement) + qualitative (sentiment, champion strength)
- **Expansion motions** — natural growth, cross-sell, multi-product, new-team expansion
- **Voice of customer into product** — how practitioner input becomes roadmap

Frameworks & Sources

- **Lincoln Murphy** — "**Customer Success**" — the modern playbook
- **Nick Mehta (Gainsight)** — CS writing and frameworks
- **Kate Leggett (Forrester)** — customer-success research
- **Jason Fried / 37signals** — "**Shape Up**" — how product orgs think about cycles
- **gtm-agents plugins**: customer-success, account-management, customer-advocacy-orchestration, customer-journey-orchestration, renewal-orchestration

Hands-On Benchmarks

- **Mutual Success Plan** — fill one out for a real prospect with outcome targets, milestones, KPI definitions, governance cadence
- **QBR dry-run** — produce the pre-read, deck, and decisions-log with a colleague
- **Health scorecard** — design a 10-metric scorecard with weights and thresholds for an AI-native SaaS
- **Expansion play design** — three plays (usage-trigger, new-team, multi-product) as 1-page specs with signals, trigger logic, owner, messaging, and success metric

📋 **Progress: Mutual Success Plan** — ☐ Drafted ☐ Reviewed ☐ Shipped · **QBR Dry-Run** — ☐ Scheduled ☐ Completed ☐ Deck archived · **Health Scorecard** — ☐ 10 metrics defined ☐ Weights justified · **Expansion Plays** — ☐ Usage-trigger ☐ New-team ☐ Multi-product

Executive Communication & Writing Craft

Writing is the multiplier on everything else. Every Clay JD values it: *"Excellent communicator who can simplify complex concepts and drive alignment from hands-on admins to the C-suite."* At the senior level, your output is **decisions that other people make differently because of what you wrote**. Unusually clear writing of good thinking will get you hired over better-credentialed candidates every time.

Core Concepts

- **Writing as thinking** — if it's unclear on the page, it's unclear in your head
- **BLUF** — Bottom Line Up Front; conclusion first, then support
- **Minto Pyramid / SCQA** — Situation, Complication, Question, Answer with grouped supporting points
- **Amazon-style narratives** — the 1-pager and the 6-pager; read-first meetings
- **Executive brief craft** — ≤200 words, TL;DR on top, decision requested, owner, next step
- **Storytelling arcs** — setup-conflict-resolution; status-quo-threat-hope
- **Presence** — what to say when you don't know; challenging without defensiveness; pausing; concision

Frameworks & Sources

- **Barbara Minto** — **"The Pyramid Principle"** — canonical
- **Jeff Bezos's shareholder letters** — the 6-pager discipline in the wild
- **Paul Graham essays** — concise tech-adjacent prose
- **William Zinsser** — **"On Writing Well"**
- **Julian Shapiro's writing guide** — free, practical
- **For spoken presence:** Toastmasters or record yourself + self-review

Hands-On Benchmarks

- **Ship one thing per week in public** — LinkedIn post, blog, video, or walkthrough. Minimum 12 weeks. Highest-leverage practice you can adopt.
- **Write a 1-page executive brief** — e.g., "Should our AI-native SaaS adopt a GTM Engineering function in H1?" Under 250 words with a decision ask.
- **Narrative rewrite** — take one LinkedIn post and rewrite it using SCQA. Compare engagement.
- **Post-event documentarian pieces** — from your Empire State pipeline; discipline them with the frameworks above.

Meta-Skills & Professional Operating System

"High agency." "Energized by ambiguity." "You don't wait for permission." These phrases appear in every Clay JD — they're not hiring vibes, they're describing learnable operating habits. The tell is whether your *average week* looks like the average week of someone already in the role you want. This domain compounds everything else.

Core Concepts & Practices

- **First-principles reasoning** — decompose to what is known to be true, rebuild from there (Feynman, Musk, Thiel)
- **Systems thinking** — see loops, not linear causes; understand leverage points (Donella Meadows)
- **Decision journals** — record every non-trivial decision with reasoning, confidence, expected outcome; review quarterly (Shane Parrish / Farnam Street)
- **Async/written-first operating** — default to doc over meeting; default to clear commit over status update
- **Weekly review ritual** — 60–90 minutes every Friday: review goals, calibrate pipeline, capture learnings
- **Deliberate practice** — isolated skill reps against a scoring rubric, not just "doing the job"
- **Public learning in public** — ship your progress, not your polished end-state
- **Energy management** — not time management: what drains you, what recharges you, cadence accordingly
- **AI-multiplier discipline** — every repetitive task gets templated, automated, or deleted

Frameworks & Sources

- **"Thinking in Systems"** — Donella Meadows
- **"Deep Work" + "Digital Minimalism"** — Cal Newport
- **"High Output Management"** — Andy Grove
- **"The Hard Thing About Hard Things"** — Ben Horowitz
- **"Working Backwards"** — Bryar & Carr — the Amazon operating system
- **Shane Parrish / Farnam Street** — mental models
- **Patrick O'Shaughnessy** — **"Invest Like the Best"** and **Sriram Krishnan** — **"The Observer Effect"** — best long-form operator interviews

Hands-On Benchmarks

- **Weekly review template** (Notion or Apple Notes) — use it for 12 weeks. Non-negotiable.
- **Learning in public cadence** — 1 public post per week on something you learned building
- **Decision journal** — log 1 decision per week for 12 weeks; review at week 12 — calibration will improve measurably
- **Protect your three highest-leverage hours per week** — delete, delegate (to AI), or defer everything else

📅 **Progress: Weekly Review** — [] Template built [] Weeks 1–4 [] Weeks 5–8 [] Weeks 9–12 • **Decision Journal** — [] Setup [] Entries logged [] Week-12 review done • **High-Leverage Hours** — [] Identified [] Protected [] Cadence locked

12-Week Sprint Plan: Three Parallel Tracks

You cannot learn nine domains sequentially without losing momentum. Run three tracks in parallel, each with a weekly rhythm. Every 12 weeks, ship one **portfolio anchor project** so good it gets pinned on LinkedIn and linked in every cold outreach.

Track	Domains	Weekly Cadence	Output
A — Build	GTM Eng + AI/Agent (3, 5)	8–12 hrs/week, concentrated	Shipped agentic workflows, Clay builds, Empire State pipeline progress
B — Study	Systems, Data, PMM, CS (2, 4, 6, 7)	4–6 hrs/week, daily reading + 1 exercise	1 written artifact per week (memo, dashboard, positioning one-pager)
C — Ship/Public	Commercial, Writing, Meta (1, 8, 9)	Embedded into daily work	1 LinkedIn post per week; weekly review Friday

Weekly Rhythm

Mon Study 2 hrs (B) + Build 2 hrs (A)	Tue Build 3 hrs (A) + 1 post draft (C)	Wed Study 2 hrs (B) + Build 2 hrs (A)	Thu Build 3 hrs (A) + Writing craft 1 hr (C)
Fri Ship artifact + Weekly review 90 min (C)	Sat/Sun Deep reading 2–3 hrs (B); rest		

📌 **Quarterly Capstone:** Every 12 weeks, ship one **portfolio anchor project** — so good it gets pinned on LinkedIn and linked in every cold outreach. This is the compounding output of all three tracks combined.

The Search as an AI-Native GTM Motion

You are not applying for jobs. You are running a full-funnel AI-native GTM motion where the product is you, the ICP is AI-native companies where full-stack GTM is a real role, and the differentiator is the method itself. Run this funnel as rigorously as you would run a \$5M ABM pursuit.

1

Market Mapping

Define ICP, build ranked 60-company target list, tier by fit

2

Positioning

Crisp positioning statement, 3–5 public proof artifacts

3

Signal Generation

Public body of work that proves skills and generates inbound

4

Prospecting

Named contacts at each target, enriched and prioritized

5

Outreach

Multi-channel, value-first — artifact-led, not resume-led

6

Evaluation → Offer → Ramp

Run discovery on yourself; 30/60/90 before the interview

The Product

You + an operating system of agents + a body of evidence. Not a resume — a living portfolio of shipped work, public writing, and demonstrated systems thinking.

The ICP

AI-native companies where full-stack GTM is a real role, not a buzzword. Series A–C, 20–250 employees, NYC or fully remote, GTM leader is a practitioner. The persona is the hiring manager, the skip-level, and the founding GTM operator — different plays, different content for each.

ICP Definition & Target Company List

Build a ranked, tiered list of ~60 target companies matching "AI-native + full-stack-GTM-compatible." Treat it exactly like an ABM target list. The goal is precision, not volume — a smaller list with deeper research outperforms a larger spray-and-pray approach every time.

ICP Criteria (Written, Not Vibes)

- **AI is core to the business** — AI-native product, or company re-architecting around AI
- **Stage: Series A–C** — larger than seed (revenue + GTM function exists); smaller than Series D+ (role shape still fluid)
- **20–250 employees** — big enough to need GTM structure, small enough to hire a 1-person full-stack function
- **NYC or fully remote** — matches your standing context
- **GTM leader is a practitioner** — not a career COO; check LinkedIn backgrounds
- **Public signal** — company values GTM engineering / AI-native GTM on their public surface

Sources

- Clay's customer roster (Intercom, Verkada, Anthropic, OpenAI, Rippling, Notion, Ramp, Mercury, Canva)
- YC W24/S24/W25/S25 AI-native batches; Work at a Startup; Interviewing.io's AI-native board
- GTMFund, Craft Ventures, Accel AI portfolios
- ICONIQ, Bessemer, a16z, Greylock, Sequoia AI-portfolio pages
- Contrary Capital, South Park Commons, Founders Fund
- Apollo saved search: AI-native + fit criteria

Tiering Framework

Tier 1 — 10 to 12 Companies Dream fits. Deserve full customized research, personalized outreach, and original work artifacts. Every touch is bespoke.	Tier 2 — ~25 Companies Strong fits. Lightly personalized approach; leverage content + network. Semi-customized outreach with real relevance signals.	Tier 3 — ~25 Companies Volume tier. Templatized application + inbound interest. Lower individual investment, broader coverage.
---	--	--

-  **Tools:** Apollo · Clay free tier · LinkedIn Sales Navigator · HubSpot (treat target companies as accounts, hiring managers as contacts) · Cowork sales:account-research skill
- Metrics:** Target list size by tier (60 total) · Warm-path coverage (% of Tier 1 with a connector) · Signal density per account (signals logged per company per month)
- Top 1% move:** Publish a version of your target list as strategic commentary — e.g., "The 10 AI-native companies most likely to define the next phase of GTM." Position as market-analyst level, not candidate level.

Named Tier-1 Anchors (from Clay Blog)

These companies explicitly run named GTM-engineering functions. Research each human by name, follow them, engage with their content, map warm paths:

- **Intercom** — Alexander DeMoulin (GTM Ops)
- **Canva** — Robert Jones (GTM AI)
- **Notion** — Evan Peters, Jalal Iftikhar, Shelley Lee, Theo Bleier
- **Anthropic** — Adam Wall (GTM Infrastructure)
- **Ramp** — Keyan Sarrafzadeh (PM, Growth)
- **Verkada** — Davide Grieco (Growth, CMO org)
- **Rippling** — Noah Adelstein (International + US Growth)

Also follow **Keith Jones** and **Rachel Hepworth** — Clay insiders who shaped the blog. Worth following for movement-level signal.

Positioning & Personal Brand Platform

A crisp, defensible positioning statement about you, supported by 3–5 public proof artifacts, so that any warm connection can explain you in one sentence and any stranger can Google you and understand you in 90 seconds.

Your Positioning One-Pager (April Dunford Template)

- **Competitive alternatives** — what hiring managers would consider instead: a traditional AE + a RevOps analyst; a GTM generalist from consulting; an internal promotion
- **Unique attributes** — 12+ years enterprise, NYC AI-event documentarian, ships agentic systems weekly, builds in public, has lived the workflow Clay and Clay-customers are automating
- **Value** — compressed hiring: 1 senior hire covers AE + ops + AI-tooling + content; you are the consultative-seller-who-built-it archetype
- **Best-fit customers** — Series A-C AI-native, 50–150 employees, GTM leader wants a generalist senior
- **Market category** — pick one: "Full-stack GTM hire," "Founding GTM Engineer + AE," "Enterprise AE with operator instincts," or "Consultative seller for AI-native GTM products"

Narrative hook: "The best sellers of this product are people who have lived the problem — lawyers selling Harvey, therapists selling Prompt Health, GTM engineers selling Clay. I'm the enterprise AE who's lived twelve years of the workflow these tools are automating, and now I'm building on the other side of it."

LinkedIn Overhaul

- **Headline:** role-language that matches your positioning, not job-title-language
- **About:** SCQA structure — Setup (what's changing in GTM), Complication (most people on one side of the AI fence), Question (what does full-stack GTM look like), Answer (how you practice it)
- **Experience:** rewrite every role into value-delivered with metrics — not responsibilities
- **Featured:** 3–5 pinned posts that are your proof artifacts

Personal Site + Surface Area Audit

- Minimal site (Framer or Vercel + one Lovable/Bolt page): positioning, body of work, current projects, contact — 1 weekend build
- Audit GitHub (make it presentable), Twitter/X if relevant, public Notion workspace
- Goal: no conflicting signal anywhere a hiring manager might land

Tools


- Cowork marketing:brand-review skill on every bio/profile
- A 5–10 min Loom walking through your work — becomes a reusable asset

📌 **Metrics:** Inbound recruiter outreach per week (leading indicator) · Profile views per week · Warm intro rate (% of target companies where you have or can manufacture a warm intro within 2 degrees) · Post engagement trend

Top 1% move: Their positioning is so specific that a hiring manager sees the one-pager and writes to them before a role opens. Optimize for fit, not breadth.


The Flywheel That Runs While You Sleep

A public body of work that proves the skill domains from Part 1, generates inbound interest from target companies, and functions as legitimate cold-outreach material. Your Empire State Events pipeline is the signal engine — every event attended becomes 1 LinkedIn post, 2–3 DMs to speakers and hosts, and 1 project idea. Over 12 weeks: ~24 events → ~24 posts + ~60 warm touches + 10–15 shipped small projects.




Daily (Micro)

A signal you observed — a tool you tried, a pattern you saw, a thread you joined. LinkedIn, 3–5 sentences. Low effort, high frequency, builds presence and pattern-recognition signal over time.



Weekly (Mid)

A real post — an event recap, a framework, a walkthrough. 150–400 words. This is your primary cadence. Use SCQA or two-thesis contrast structure. Max 1/week for two-thesis posts — needs two briefs.



Monthly (Macro)

A long-form — a strategic narrative, a build write-up, a synthesis essay. 1,200–2,000 words. This is where you become the person hiring managers think of when they think about the topic. PMM positioning applied to yourself.

The Build-Writeup Template (Signals Rigor)

Every shipped-project post follows this skeleton — six sections, 250–400 words, 1 diagram. This is exactly the shape of build snippets in the Clay blog. Imitating the shape signals fluency:

1. The problem in one sentence
2. The rung of the Clay maturity model this build serves
3. The signal/data source
4. The workflow shape in 3–5 bullets
5. The measurement — meetings booked (actual or hypothesized) + hours saved vs. manual baseline
6. What you'd build next

Target-Aware Content

~20% of your content should be directly relevant to a Tier 1 company's current moment — a new product launch, a recent raise, a new content piece from them. You're not stalking; you're engaging as a peer.

Tools

- Empire State pipeline (the signal engine)
- Cowork marketing:draft-content and marketing:content-creation as draft partners
- content-correspondent and post-event-content-generator skills (already installed)
- LinkedIn Creator Mode + Featured section

📌 **Metrics:** Weekly posts published (hit rate ≥95%) · Impressions + engagement rate trend (arcs over 8–12 weeks, not week-to-week) · Warm inbound per week (DMs from new targets) · Content → meeting conversion · For any build-writeup post: meetings booked + hours saved (adopt Clay's measurement vocabulary — this is the language hiring managers at your Tier 1 targets use natively)

Top 1% move: Publishes a recurring, ownable format — e.g., a monthly "State of GTM Engineering" essay. Hiring managers read it because they want the information, and you become the person they think of when they think about the topic. Your documentarian angle from the Empire State pipeline already *is* an ownable format — productize it.

Prospecting & Outreach: Multi-Channel, Value-First

For every Tier 1 and Tier 2 company, identify 3-5 people who could credibly hire you or influence your being hired. Then run at least one real conversation with a decision-maker or influencer that either becomes a formal process, produces a referral, or earns a committed "keep you on my radar."

Stage 4 — Prospecting: Who to Reach

For each target, identify and enrich:

- **Hiring manager** — VP Sales / CRO / Head of GTM / founding GTM operator
- **Skip-level** — CEO at Series A-B; CRO/COO at Series C
- **Peer in the function** — someone doing a slice of the role today
- **Recruiter** — in-house if they have one; external talent partner otherwise
- **Connector path** — someone in your network 1-2 degrees from any of the above

Enrich with Apollo, LinkedIn Sales Nav, and Clay. Fields: LinkedIn URL, email, role tenure, last public content, recent activity. Maintain the target→contact map in HubSpot.

Stage 5 — Outreach Paths

- **Warm intro path:** ask your connector with a 3-sentence forwardable context — what you're looking for, why this company, what you'd offer. Make the ask effortless (one Slack/DM, one-click forward).
- **Content-engagement path:** thoughtful comment on their content for 2-3 weeks, then DM referencing a specific exchange. Do not DM with "would love to chat."
- **Cold artifact-gift path:** send a personalized artifact — a competitive teardown, a workflow you built, a Loom walking through a play. 3-5 min of video is disproportionately effective for senior outreach.

Message Craft

- **First touch:** relevance signal (why you, why now, why them), offer (a useful artifact or specific observation), ask (small — 20 min, not "pick your brain")
- **Follow-up 1 (5 business days):** new artifact or new angle; never "just bumping this up"
- **Follow-up 2 (10 business days):** a share-and-close; "sharing this in case it's useful, no need to respond"

Discovery Call — Run It Like a Sales Call

- Their current-state pain (hiring gap? GTM misfit? motion change?)
- Impact (what does not solving it cost them)
- Decision process (who else is involved, what's the timeline)
- Next step (a clear, time-bound follow-up — never leave without it named)

Post-Conversation Artifact

Within 24 hours, send a 1-page "what I heard + what I'd do in the first 90 days" memo. This single act converts interest to offers at a rate that defies most expectations — it signals you operate at a strategic level before you're hired.

 **Tools:** Gmail + HubSpot for CRM logging · Loom or Claude-driven video (via ElevenLabs) for video touches · gtm-agents sales-prospecting/* skills for sequence design · Cowork sales:draft-outreach skill

Metrics: Outreach sent per week (target: 10-15 tightly-personalized touches) · Reply rate (target: 25-40% warm, 10-20% cold-with-artifact) · Conversations booked per week (target: 2-4 at steady state) · "Process advanced or recycled" per conversation (want: ≥60%)

Top 1% move: Sends a Loom before the interview offering a teardown of the company's current state and a 30/60/90. About half of hiring managers skip the rest of the process when they get one of these from someone whose credentials obviously match. An hour of your time vs. 8 hours of interview process avoided. This is the proof that you operate like a Forward Deployed GTME before the offer.